

2b: PyTorch

Week 3: Overview

This week, we will look at the basic structure and components of a typical PyTorch program, and run some simple examples. We will also learn how to analyze the hidden unit dynamics of neural networks.

Weekly learning outcomes

By the end of this module, you will be able to:

- code simple PyTorch operations
- analyze the geometry of hidden unit activations in neural networks

PyTorch

The following code fragments illustrate the typical structure of a PyTorch program, with further details and various options for each component.

Typical Structure of a PyTorch Program

```
# create neural network according to model specification
net = MyModel().to(device) # CPU or GPU

# prepare to load the training and test data
train_loader = torch.utils.data.DataLoader(...)
test_loader = torch.utils.data.DataLoader(...)

# choose between SGD, Adam or other optimizer
optimizer = torch.optim.SGD(net.parameters,...)

# enter the training loop
for epoch in range(1, epochs):
    train(params, net, device, train_loader, optimizer)
    # periodically evaluate the network on the test data
    if epoch % 10 == 0:
        test(params, net, device, test_loader)
```

Defining a model

```
class MyModel(torch.nn.Module):

    def __init__(self):
        super(MyModel, self).__init__()
        # define structure of the network here

    def forward(self, input):
        # apply network and return output
```

Defining a Custom Model

This code defines a module for computing a function of the form $(x, y) \mapsto Ax \log(y) + By^2$

```
import torch.nn as nn
```

```

class MyModel(nn.Module):

    def __init__(self):
        super(MyModel, self).__init__()
        self.A = nn.Parameter(torch.randn((1),requires_grad=True))
        self.B = nn.Parameter(torch.randn((1),requires_grad=True))

    def forward(self, input):
        output = self.A * input[:,0] * torch.log(input[:,1]) \
            + self.B * input[:,1] * input[:,1]
        return output

```

Building a Net from Individual Components

```

class MyModel(torch.nn.Module):

    def __init__(self):
        super(MyModel, self).__init__()
        self.in_to_hid = torch.nn.Linear(2,2)
        self.hid_to_out = torch.nn.Linear(2,1)

    def forward(self, input):
        hid_sum = self.in_to_hid(input)
        hidden = torch.tanh(hid_sum)
        out_sum = self.hid_to_out(hidden)
        output = torch.sigmoid(out_sum)
        return output

```

Defining a Sequential Network

```

class MyModel(torch.nn.Module):

    def __init__(self, num_input, num_hid, num_out):
        super(MyModel, self).__init__()
        self.main = nn.Sequential(
            nn.Linear(num_input, num_hid),
            nn.Tanh(),
            nn.Linear(num_hid, num_out),
            nn.Sigmoid()
        )

    def forward(self, input):
        output = self.main(input)
        return output

```

Sequential Components

Network Layers:

- `nn.Linear()`
- `nn.Conv2d()` [Week 4]

Intermediate Operators:

- `nn.Dropout()`
- `nn.BatchNorm()` [Week 4]

Activation Functions:

- `nn.Sigmoid()`
- `nn.Tanh()`
- `nn.ReLU()` [Week 3]

Declaring Data Explicitly

```
import torch.utils.data

# input and target values for the XOR task
input = torch.Tensor([[0,0],[0,1],[1,0],[1,1]])
target = torch.Tensor([[0],[1],[1],[0]])

xdata = torch.utils.data.TensorDataset(input,target)
train_loader = torch.utils.data.DataLoader(xdata,batch_size=4)
```

Loading Data from a .csv File

```
import pandas as pd

df = pd.read_csv("sonar.all-data.csv")
df = df.replace('R',0)
df = df.replace('M',1)
data = torch.tensor(df.values,dtype=torch.float32)
num_input = data.shape[1] - 1
input = data[:,0:num_input]
target = data[:,num_input:num_input+1]
dataset = torch.utils.data.TensorDataset(input,target)
```

Custom Datasets

```
from data import ImageFolder
# load images from a specified directory
dataset = ImageFolder(folder, transform)
```

```
import torchvision.datasets as dsets
# download popular image datasets remotely
mnistset = dsets.MNIST(...)
cifarset = dsets.CIFAR10(...)
celebset = dsets.CelebA(...)
```

Choosing an Optimizer

```
# SGD stands for “Stochastic Gradient Descent”
optimizer = torch.optim.SGD( net.parameters(),
    lr=0.01, momentum=0.9,
    weight_decay=0.0001)

# Adam = Adaptive Moment Estimation (good for deep networks)
optimizer = torch.optim.Adam(net.parameters(),eps=0.000001,
    lr=0.01, betas=(0.5,0.999),
    weight_decay=0.0001)
```

Training

```
def train(args, net, device, train_loader, optimizer):

    for batch_idx, (data,target) in enumerate(train_loader):
        optimizer.zero_grad() # zero the gradients
        output = net(data)    # apply network
        loss = ...            # compute loss function
        loss.backward()       # update gradients
        optimizer.step()      # update weights
```

Loss Functions

```
loss = torch.sum((output-target)*(output-target))
loss = F.nll_loss(output,target)                # [Week 3]
loss = F.binary_cross_entropy(output,target)    # [Week 3]
loss = F.softmax(output,dim=1)                  # [Week 3]
loss = F.log_softmax(output,dim=1)              # [Week 3]
```

Testing

```
def test(args, net, device, test_loader):
    with torch.no_grad(): # suppress updating of gradients
```

```
net.eval()          # toggle dropout, batch norm [Week 3]
test_loss = 0
for data, target in test_loader:
    output = model(data)
    test_loss += ...
print(test_loss)
net.train()         # toggle dropout, batch norm back again
```

Computational Graphs

PyTorch automatically builds a computational graph, enabling it to backpropagate derivatives.

Every parameter includes `.data` and `.grad` components, for example:

```
A.data
```

```
A.grad
```

`optimizer.zero_grad()` sets all `.grad` components to zero.

`loss.backward()` updates the `.grad` component of all Parameters by backpropagating gradients through the computational graph.

`optimizer.step()` updates the `.data` components.

Controlling the Computational Graph


If we need to stop the gradients from being backpropagated through a certain variable (or expression) `A`, we can exclude it from the computational graph by using:

```
A.detach()
```

By default, `loss.backward()` discards the computational graph after computing the gradients.

If needed, we can force it to keep the computational graph by calling it this way:

```
loss.backward(retain_graph=True)
```



Exercise: Running PyTorch

The following program solves the simplest possible machine learning task:

solve $f(x) = wx$ such that $f(1) = 1$

```
import torch
import torch.utils.data
import numpy as np

lr = 1.9 # learning rate
mom = 0.0 # momentum

class MyModel(torch.nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        self.w = torch.nn.Parameter(torch.zeros((1), requires_grad=True))
    def forward(self, input):
        output = self.w * input
        return(output)

device = 'cpu'

input = torch.Tensor([[1]])
target = torch.Tensor([[1]])

slope_dataset = torch.utils.data.TensorDataset(input,target)
train_loader = torch.utils.data.DataLoader(slope_dataset,batch_size=1)

# create neural network according to model specification
net = MyModel().to(device) # CPU or GPU

# choose between SGD, Adam or other optimizer
optimizer = torch.optim.SGD(net.parameters(),lr=lr,momentum=mom)

epochs = 1000

for epoch in range(1, epochs):
    for batch_id, (data,target) in enumerate(train_loader):
        optimizer.zero_grad() # zero the gradients
        output = net(data) # apply network
        loss = 0.5*torch.mean((output-target)*(output-target))
        print('Ep%3d: zero_grad(): w.data=%7.4f loss=%7.4f' \
              % (epoch, net.w.data, loss))
        loss.backward() # compute gradients
        optimizer.step() # update weights
        print('step(): w.grad=%7.4f w.data=%7.4f' \
              % (net.w.grad, net.w.data))
        if loss < 0.000000001 or np.isnan(loss.data):
            exit(0)
```


Question 1

Change the learning rate `lr` to each of the following values by editing line 5 in the above code.

0.01, 0.1, 0.5, 1.0, 1.5, 1.9, 2.0, 2.1

Try running the code and describe what happens for each value of `lr`, in terms of the success and speed of the algorithm.

No response

Question 2

Now keep the learning rate at `1.9`, but try each of the following values for momentum by changing the value of `mom` on line 6.

0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9

For which value of momentum is the task solved in the fewest epochs?

What happens when the momentum is `1.0`? What happens when it is `1.1`?

No response

Exercise: XOR with PyTorch

This program trains a two-layer neural network on the famous XOR task.

```
import torch
import torch.utils.data
import torch.nn.functional as F

lr = 0.1
mom = 0.0
init = 1.0

class MyModel(torch.nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        # define structure of the network here
        self.in_hid = torch.nn.Linear(2,2)
        self.hid_out = torch.nn.Linear(2,1)
    def forward(self, input):
        # apply network and return output
        hid_sum = self.in_hid(input)
        hidden = torch.tanh(hid_sum)
        out_sum = self.hid_out(hidden)
        output = torch.sigmoid(out_sum)
        return(output)

device = 'cpu'

input = torch.Tensor([[0,0],[0,1],[1,0],[1,1]])
target = torch.Tensor([[0],[1],[1],[0]])

xor_dataset = torch.utils.data.TensorDataset(input,target)
train_loader = torch.utils.data.DataLoader(xor_dataset,batch_size=4)

# create neural network according to model specification
net = MyModel().to(device) # CPU or GPU

# initialize weight values
net.in_hid.weight.data.normal_(0,init)
net.hid_out.weight.data.normal_(0,init)

# choose between SGD, Adam or other optimizer
optimizer = torch.optim.SGD(net.parameters(),lr=lr,momentum=mom)

epochs = 10000

for epoch in range(1, epochs):
    #train(net, device, train_loader, optimizer)
    for batch_id, (data,target) in enumerate(train_loader):
        optimizer.zero_grad() # zero the gradients
```

```
output = net(data)    # apply network
loss = F.binary_cross_entropy(output,target)
loss.backward()        # compute gradients
optimizer.step()       # update weights
if epoch % 100 == 0:
    print('ep%3d: loss = %7.4f' % (epoch, loss.item()))
if loss < 0.01:
    print("Global Mininum")
    exit(0)
print("Local Minimum")
```

Question 1

Run the above code ten times. For how many runs does it reach the Global Minimum? For how many runs does it reach a Local Minimum?

No response

Question 2

Keeping the learning rate fixed at `0.1`, adjust the values of momentum (`mom`) on line 6 and initial weight size (`init`) on line 7 to see if you can find values for which the code converges relatively quickly to the Global Minimum on virtually every run.

No response

Coding Exercise: Basic PyTorch Operations



Objective

The **Tensor** is a fundamental structure in PyTorch which is very similar to an array or matrix. Tensors are used to encode the inputs and outputs of a model, as well as the model's parameters. In this exercise, you will learn how to implement basic tensor operations.

Instructions

Before starting the exercise, please go through the tutorial about tensors from the PyTorch website.

https://pytorch.org/tutorials/beginner/blitz/tensor_tutorial.html#sphx-glr-beginner-blitz-tensor-tutorial-py

For some of the exercises, the `torch.Tensor` documentation should be very helpful.

<https://pytorch.org/docs/stable/tensors.html>

Week 2 Thursday video



Week 3 Tutorial

[Week 3 Tutorial Questions](#)